

Diffusion Lens: Interpreting Text Encoders in Text-to-Image Pipelines

Michael Toker Hadas Orgad Mor Ventura Dana Arad Yonatan Belinkov

Technion – Israel Institute of Technology

{tok, orgad.hadas, mor.ventura, danaarad, belinkov}@campus.technion.ac.il

Abstract

Text-to-image diffusion models (T2I) use a latent representation of a text prompt to guide the image generation process. However, the process by which the encoder produces the text representation is unknown. We propose the DIFFUSION LENS, a method for analyzing the text encoder of T2I models by generating images from its intermediate representations. Using the DIFFUSION LENS, we perform an extensive analysis of two recent T2I models. Exploring compound prompts, we find that complex scenes describing multiple objects are composed progressively and more slowly compared to simple scenes; Exploring knowledge retrieval, we find that representation of uncommon concepts require further computation compared to common concepts, and that knowledge retrieval is gradual across layers. Overall, our findings provide valuable insights into the text encoder component in T2I pipelines.¹

1 Introduction

The text-to-image (T2I) diffusion pipeline comprises two main elements: the text encoder and the diffusion model. The text encoder converts a textual prompt into a latent representation, while the diffusion model utilizes this representation to generate the corresponding image. Several recent studies have delved into the internal workings of the diffusion model and the cross-attention mechanism that connects the two components (Tang et al., 2023; Hertz et al., 2023; Orgad et al., 2023; Chefer et al., 2023a). Yet, while the text encoder is a key component of the pipeline with a large effect on image quality and text-image alignment (Saharia et al., 2022), its internal mechanisms remain unexplored. Moreover, while there is a wide range of work that has analyzed general language model internals (Belinkov and Glass, 2019; Rogers et al.,

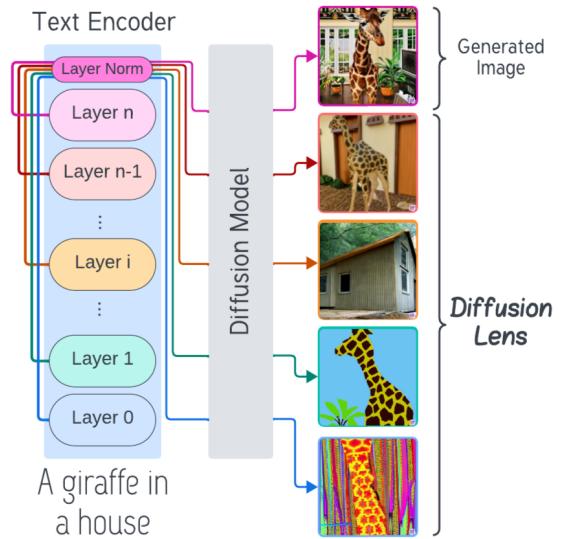


Figure 1: Visualization of the text encoder’s intermediate representations using the DIFFUSION LENS. At each layer of the text encoder (in blue), the DIFFUSION LENS takes the full hidden state, passes it through the final layer norm, and feeds it into the diffusion model.

2020; Madsen et al., 2022), these methods are not suitable for exploring fine-grained visual features.

We propose the DIFFUSION LENS, a method for analyzing the inner mechanism of the text encoder. The DIFFUSION LENS uses intermediate representations of the prompt from various layers of the text encoder to guide the diffusion process, resulting in images that are clear, consistent, and easy to understand for most layers (see Figure 1). Notably, the DIFFUSION LENS relies solely on the pre-trained weights of the model and does not depend on any external modules.

We employ the DIFFUSION LENS to examine the computational process of the text encoder in two popular T2I models: Stable Diffusion (Rombach et al., 2022) and Deep Floyd (StabilityAI, 2023). Our investigation focuses on two main analyses: the model’s capability of conceptual combination

¹Code and data are available on the project webpage tokeron.github.io/DiffusionLensWeb.

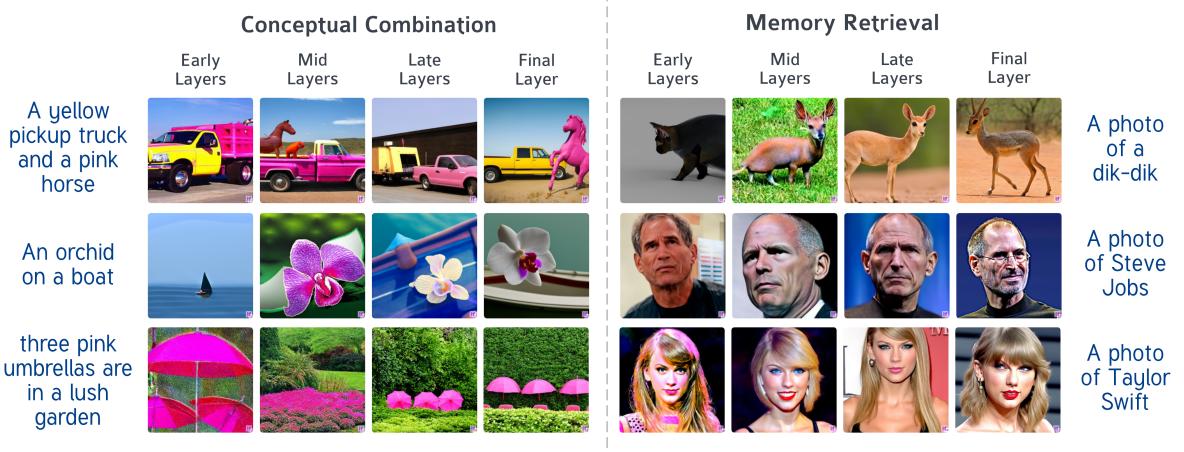


Figure 2: Insights gained from using DIFFUSION LENS. **Conceptual Combination** (left): early layers often act as a “bag of concepts”, lacking relational information which emerges in later layers. **Memory Retrieval** (right): uncommon concepts gradually evolve over layers, taking longer to generate compared to common concepts.

and its memory retrieval process. For each analysis, we either construct a tailored dataset to isolate a specific phenomenon or utilize naturally occurring human-written image captions.

Our analysis of conceptual combination reveals various insights: (1) Complex prompts such as “A yellow pickup truck and a pink horse” require more computation to achieve a faithful representation compared to simpler prompts such as “A green cat”. (2) Complex representations are built gradually: as illustrated in Figure 22 (Left), images generated from early layer representations typically encode concepts separately or together without capturing their correct relationship, resembling more of a “bag of concepts”. Images from subsequent layers encode the relationships more accurately. (3) The order in which objects emerge during computation is influenced by either their linear or syntactic precedence in the sentence. Here, we find a difference between the two examined models: Deep Floyd’s text encoder, T5 (Raffel et al., 2020), shows a greater sensitivity to syntactic structure, while Stable Diffusion’s text encoder, CLIP (Radford et al., 2021), tends to reflect linear order.

Next, we investigate memory retrieval, and uncover several key findings: (1) Common concepts, such as “Kangaroo”, emerge in early layers while less common ones, such as the animal “Dik-dik”, gradually emerge across the layers, with the most accurate representations predominantly occurring in the upper layers, as illustrated in Figure 22 (Right, top). (2) Fine details, like human facial features, materialize at later layers, as shown in Fig-

ure 22, with the prompt “A photo of Steve Jobs”. (3) Knowledge retrieval is gradual, unfolding as computation progresses. This observation diverges from prior research on knowledge encoding in language models which characterizes knowledge as a localized attribute encoded in specific layers (Geva et al., 2022; Meng et al., 2022; Arad et al., 2023). (4) Notably, there are discernible differences in memory retrieval patterns between the two text encoders: Deep Floyd’s T5 memory retrieval exhibits a more incremental behavior compared to Stable Diffusion’s CLIP. The disparities uncovered through our analyses suggest that factors such as architecture, pretraining objectives, or data may influence the encoding of knowledge or language representation within the models.

Our contributions are summarized as follows:

- We develop the DIFFUSION LENS, a new intrinsic method for analyzing the intermediate states of the text encoder within T2I pipelines.
- We conduct thorough experiments that reveal insights on the computational mechanisms of text encoders in the T2I pipeline. Our findings shed light on how factors such as complexity, frequency, and syntactic structure impact the encoding process.

2 Diffusion Lens

Preliminaries. Current text-to-image diffusion models comprise two main components (Saharia et al., 2022; Ramesh et al., 2022): a language model used as a text encoder that takes the textual prompt as input and produces latent representations; and

a diffusion model that is conditioned on the representations from the text encoder and generates an image from an initial input noise.

The language model in the T2I pipeline is typically a transformer model. Transformer models consist of a chain of transformer blocks, each composed of three sub-blocks: attention, multi-layer perceptron, and layer norm (Vaswani et al., 2017).

We denote the transformer block at layer l as F_l . The input to the model is a sequence of T word embeddings, denoted as $\mathbf{h}^0 = [h_1^0, \dots, h_T^0]$. Then, the output of the transformer block at layer l is a sequence of hidden states \mathbf{h}^{l+1} :

$$\mathbf{h}^{l+1} = F_l(\mathbf{h}^l) \quad (1)$$

The output representations of the last block, L , go through a final layer norm, denoted as ln_f . Then, they condition the image generation process through cross-attention layers, resulting in an image I . We abstract this process as:

$$I = \text{Diff}(ln_f(\mathbf{h}^L)) \quad (2)$$

Diffusion Lens. In a T2I pipeline with a text encoder of L layers, for layer $l < L$, we process the output of block l , including padding tokens, through the final layer norm. We condition the diffusion process on this output, as illustrated in Figure 1. Namely, we generate an image I from an intermediate layer l as follows:

$$I = \text{Diff}(ln_f(\mathbf{h}^l)) \quad (3)$$

The final layer norm is a crucial step in generating coherent images (see Appendix A.3). It projects the representations into the cross-attention embedding space without the caveat of adding new information to the representation, as may happen with learned projections. This process generates an image representing the intermediate state of the text-encoder as interpreted by the diffusion model.

3 Experimental Setup

Models. The experiments are performed on Stable Diffusion 2.1 (denoted *SD*, Rombach et al., 2022) and Deep Floyd (denoted *DF*, StabilityAI, 2023). SD is an open-source implementation of latent diffusion (Rombach et al., 2022), with OpenCLIP-ViT/H (Ilharco et al., 2021) as the text-encoder. DF is another open-source implementation of latent diffusion inspired by Saharia et al. (2022), with a frozen T5-XXL (Raffel et al., 2020) as the text encoder. We usually only report the results on DF, unless there is a difference between

the models, which we then discuss. The full results on SD are given in Appendix E.

Data. Depending on the specific experiment, we either curate prompt templates and automatically generate a list of prompts from a collected list of concepts we are interested in investigating, or use a list of natural, handwritten prompts from COCO (Lin et al., 2015). The data for each experiment is detailed in the next sections. With each prompt, we generate images that are conditioned on representations from every fourth layer in the model, which serves as a representative subset. This results in 7 images for DF (which has 25 layers in total) and 6 images for SD (which has 24). We generate each prompt using four seeds.

Evaluation. In every experiment we ask questions about the images at every layer, e.g., “Does the prompt correspond to the generated image”; or, if there are two objects in the prompt, “Does object A appear in the generated image?”. We describe the questions in detail for every experiment below. To analyze the representation building process of successful generations, we report our main findings on cases where all the images from the last layer align with the prompt. We separately analyze model failures in Section 6.

We annotated the generated images using both human annotators and GPT-4V (OpenAI, 2023). For the human evaluation, we collected answers to the questions by ten human annotators, with 10% overlap to measure inter-annotator agreement. We found a high agreement between GPT-4V and the humans (Table 2, App. B). We provide the main results based on the human annotations; however, our results support the use of automatic annotation to allow larger scale and reduced cost. Overall, we collected answers to roughly 66,560 questions, 37% of them by GPT-4V. For full details on the annotation process, inter-annotator agreement, and integration with GPT-4V, refer to Appendix B.

4 Conceptual Combination

T2I diffusion models are popular for their ability to generalize beyond their training data, creating composite concepts (Ramesh et al., 2022). Conceptual combination is the cognitive process by which at least two existing basic concepts are combined to generate a new higher-order, composite concept (ling Wu and Barsalou, 2009). Conceptual combination is at the core of knowledge representation, since it asks how the meaning of a complex phrase

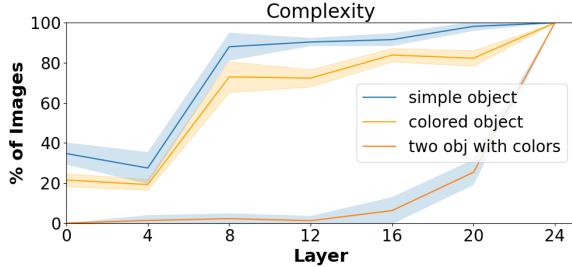


Figure 3: Percentages of prompt-matching images across various layers. As prompts become more complex, DIFFUSION LENS has to utilize more layers to extract a correct image.

connects to its component parts (Hampton, 2013), e.g., “A cat in a box”. This section uses the DIFFUSION LENS to trace the process by which the text encoder creates composite concepts.

4.1 Building complex scenarios

This study investigates the text encoder’s ability to combine concepts at varying levels of complexity. We utilize COCO classes (Lin et al., 2015) as a diverse set of prompts with readily identifiable visual meanings. Each experiment commences with a simple list of objects as prompts, progressively increasing in complexity as outlined subsequently.

Colors and conjunction. We compile three lists of prompts: (1) objects (e.g., “a dog”); (2) objects with color description (“a red dog”); and (3) two objects with colors (“a red dog and a white cat”). To investigate how conceptual combination emerges through the layers, we annotated a random sample of 80 prompts,² asking the following questions for each layer: (a) Does object X appear in the image? (b) Does color X appear in the image? (c) Does object X appear in the correct color? X is either the 1st or the 2nd object, for a total of 6 questions.

Physical relations. We compile two lists of prompts: (1) objects, (2) prompts describing two objects and a preposition: either “in” or “on”. For example, “A cat in a box”. We sample 40 prompts. We ask three questions: (a-b) Does object X appear in the image? and (c) Is object A in/on object B?

Results

The simpler the concept, the earlier it emerges. Figure 3 shows the percentage of images that correctly generated the concepts for each category: an

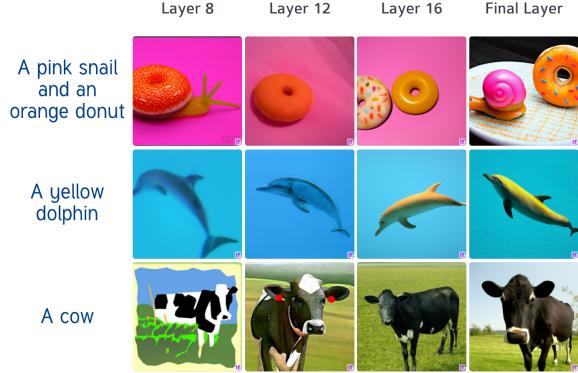


Figure 4: Complex prompts take more computation blocks to emerge.

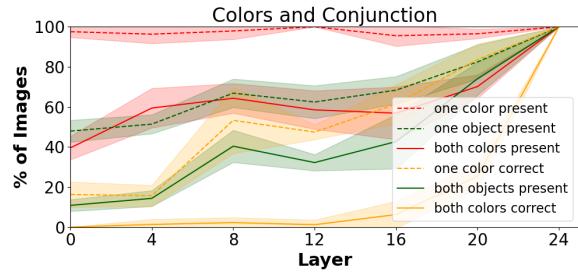


Figure 5: The proportion of images where either the object, the colors, or both were present, and where either the objects or the colors were accurately represented.

object, an object and a color, and two colored objects. Prompts describing a single object emerge the earliest, between layers 4 and 16, while prompts containing a color descriptors emerge in layers 16–20. Conjunction prompts emerge last, around layers 20–24. As demonstrated in Figure 4, “A cow” is fully represented by layer 8, while “A yellow dolphin” does not correctly form until layer 16. Lastly, “A pink snail and an orange donut” only fully forms at much later layers, correctly matching the objects and colors at the final layer, 24.

Complex representations are constructed gradually. We continue with the complex prompts of two colored objects. Figure 5 aggregates the answers to illustrate the behavior of either or both objects appearing in intermediate layers. Colors often emerge first, with both colors often emerging in early layers in DF (in SD, the two objects appear before two colors). A single object is also gradually represented in layers 4–12. Notably, while the colors and one of the objects appear, the object is not necessarily generated in the correct color. This can be seen in the first example in Figure 6: While a raccoon and a rocket do appear, and the image

²In this experiment, human annotators annotated 40 prompts and GPT4-V annotated an additional 40.

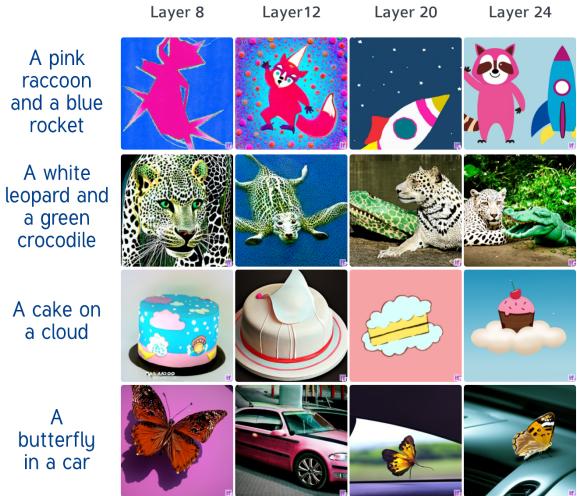


Figure 6: Complex representations are constructed gradually. In some cases, objects are mixed in early representations. In other cases, only one of the objects appear in early representations.

| Model | Antecedent first | | Antecedent second | |
|-----------|----------------------|----------------------|----------------------|----------------------|
| | 1 st noun | 2 nd noun | 1 st noun | 2 nd noun |
| DF (T5) | 50.8% | 33.87% | 35.50% | 51.60% |
| SD (Clip) | 58.4% | 23.80% | 54.90% | 27.90% |

Table 1: The percentage of prompts in each group where the antecedent noun (either the first or the second noun mentioned) appeared earlier.

contains both blue and pink elements, the rocket is not blue until the final layer. In some cases, we observe a mixture of concepts in early layers, as seen in the second example of Figure 6. Similarly, the bottom two examples in Figure 6 show prompts composing two objects and a proposition. As with colors, we observe that individual objects appear in early layers but the correct relation emerges much later. For example, “A cake on a cloud” generates images of both a cake and a cloud, with different relations; at layer 8 the cake is decorated with a cloud and in layer 20 the clouds are depicted as frosting. The correct relation is only generated at the final layer. We provide a more detailed discussion of prepositions in Appendix A.1. The patterns we see in these prompts suggest that the early representations of the text encoder behave like a “bag of concepts”, with a representation for each concept but no clear relations between them.

4.2 Syntactic dependencies

To investigate the order in which different objects emerge, we focus on the association between syn-

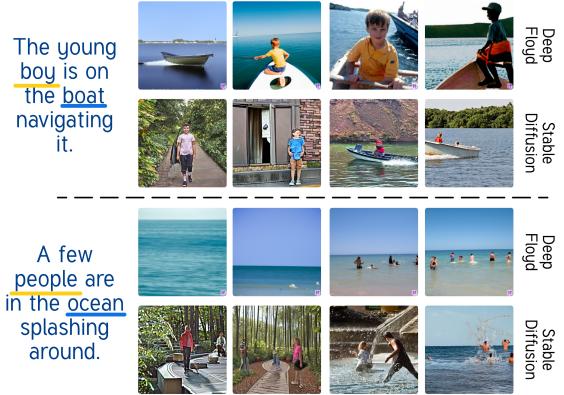


Figure 7: Difference between DF and SD models. The antecedent is marked in yellow and the descendant is marked with blue. In many cases, the antecedent appears in an earlier layer in DF, while the first noun tends to appear first in SD, regardless of its syntactic role.

tactic depth and the appearance order of nouns. Specifically, we explore whether, in a dependency path where noun A precedes noun B, noun A appears at earlier layers through DIFFUSION LENS. Using 63K prompts from COCO that we parsed with Stanza (Qi et al., 2020), we filtered for instances with two nouns per prompt and analyzed the dependency relations between the nouns. We categorized the data based on the linear position of the antecedent and generated images with 40 random samples from each group. For each generation and intermediate layer, and each object X, we queried whether object X appears in the image.

Results. First, we sometimes observe a “race” between the nouns: in 11.9% of the cases in DF, the object that appears in an earlier layer disappears at a later layer, while the other object takes dominance. See Appendix A.2 for examples.

Second, Table 1 presents information on the order of generation for both models, revealing that the sequence in which objects emerge during the computation process is determined by either their linear or their syntactic precedence, *depending on the particular text encoder*. In DF’s T5 text encoder, slightly over half of the instances feature the antecedent appearing at an earlier layer than the descendant, with a smaller fraction showing the opposite, and the rest indicating simultaneous appearances. This holds true regardless of linear order. Conversely, in SD’s CLIP, the first noun tends to appear before the second more frequently, irrespective of the syntactic role. See Figure 7 for

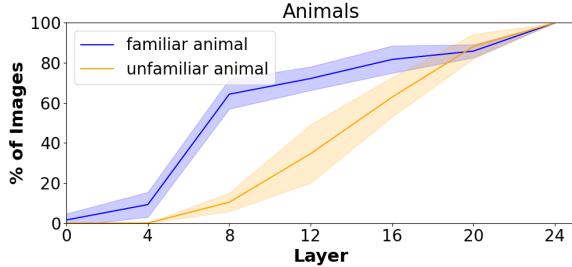


Figure 8: Common vs. uncommon animals across layers. Common animals emerge at much earlier layers.

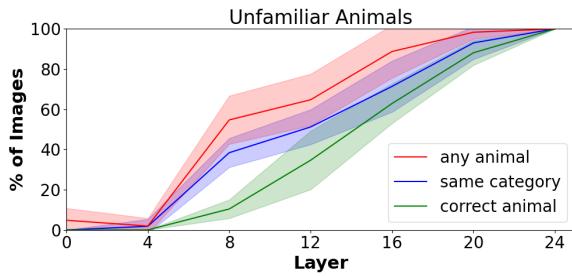


Figure 9: Subset of layers encoding different features in the process of uncommon animal generation.

a qualitative example of this case.

While the two models differ in multiple respects (architecture, pretraining data, training objective, and more), it is intriguing to observe that T5, trained on a language modeling objective, demonstrates a greater awareness of syntactic structure compared to CLIP – a model trained to align pairs of prompts and images without a specific language modeling objective. This discrepancy points to a possible impact of training objectives on the models’ representation building process.

5 Memory Retrieval

Text-to-image diffusion models are able to retrieve information of many concepts (Ramesh et al., 2022), encompassing entities like notable individuals, animals, and more. Memory retrieval—the recall of stored information— involves a constructive process rooted in the interactive dynamics between memory trace features and retrieval cue characteristics (Smelser et al., 2001). In this section, we leverage the DIFFUSION LENS to scrutinize the memory retrieval mechanism in the text encoder.

5.1 Common and Uncommon Concepts

We investigate whether there is a difference in the generation process for prompts describing common and uncommon concepts, using a list of common



Figure 10: Incremental progression in DF versus early knowledge representation in SD.

and uncommon *animals*.³ Commonality in this context does not refer to the commonality of an animal in the world, but rather to its commonality in the training data. As a proxy to measure commonality in the training data, we utilized the average daily view statistics of Wikipedia pages from October 2022 to October 2023. An animal was deemed “common” if it had an average of 1500 visits per day on its Wikipedia page (e.g., kangaroo), while one having fewer than 800 visits per day was deemed “uncommon”. We verified this distinction by examining the frequencies of species names in the LAION2B-en dataset (Schuhmann et al., 2022), extracted by Samuel et al. (2024), and found that the frequency of common species was greater than that of uncommon species with statistical significance (Appendix C).

Since the model may have seen the uncommon animals less frequently during training, their generation may take longer. We annotate each image by asking if the specific animal in the prompt appears in the generated image.

Results. As summarized in Figure 8, *common concepts emerge early*, as early as layer 8 out of 24. In contrast, *uncommon concepts gradually become apparent across the layers*, with accurate images generated primarily at the top layers.

5.2 Gradual Retrieval of Knowledge

To delve deeper into the knowledge retrieval process, we pose additional questions for uncommon animal: (a) Is there an animal in the image? (b)

³We use animal species as we found that models can correctly generate images of uncommon species, unlike uncommon objects and celebrities.



Figure 11: Intricate details are refined gradually.

Does the image feature an X? where X is the informal “category”⁴ of the animal, such as “mammal” and “bird”. (c) Does the image depict the exact animal in the prompt?

Results. Figure 9 illustrates *incremental knowledge extraction*, beginning with a general animal, progressing to a more specific animal within the same category, and reaching a representation of the particular animal mentioned in the prompt.

Though the plot for SD reveals a similar pattern (Appendix E), qualitative analysis reveals *distinct knowledge retrieval patterns between the two models*: In the case of DF’s T5, knowledge retrieval is gradual, unfolding as computation progresses (Figure 10). Layers generate animal, mammal, and ultimately construct a representation of the specific animal. However, SD’s text encoder, Clip, does not display a similar progression of retrieval. The model establishes the representation less gradually: The first layer with a meaningful image already closely resembles the final animal, with subsequent layers mainly refining its characteristics. These differences echo the syntactic findings in Section 4.2. They suggest that pretraining objectives, data, or model architecture might influence information organization, leading to distinct memory retrieval patterns.

5.3 Gradual refinement of features

As the computation progresses, both accuracy and realistic representation significantly improve with refining details at each step. This progression is evident in Figure 11 (top row), as seen in the gradual refinement of the “Tarsier” image. A similar

⁴We chose to use an informal taxonomy as the animal kingdom taxonomy is a complex subject under research and debate, and its terms are not common to the general population and, hence, likely less present in the T2I training data.

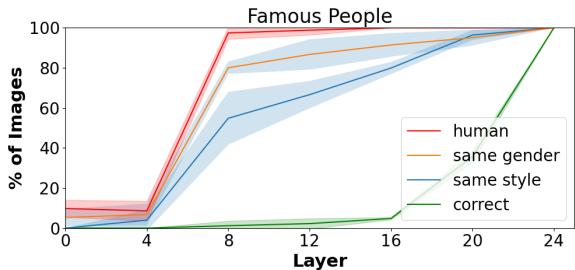


Figure 12: The distribution of feature granularity across layers in generated images.

trend occurs in the representation construction of human subjects, with facial features undergoing refinement for a more faithful portrayal (Figure 11, rows 2+3). To systematically assess this phenomenon, we compiled a list of 30 celebrities, using DIFFUSION LENS to generate images from intermediate representations in the text encoder. For each prompt and generated image, we ask: (a) Is there a person in the image? (b) Does the person align with the celebrity’s (self-identified) gender? (c) Does the person exhibit the celebrity’s style (hair, clothing, etc.)? (d) Is the individual in the image distinctly recognizable as the specified celebrity based on facial features?

Results. Figure 12 quantifies the *step-by-step construction of the representation*, culminating in its maximum resemblance to the celebrity. The integration of distinct features follows a hierarchical pattern, progressing from broad characteristics (such as the overall human form) to finer details (specifically, facial features), which become evident only in the final layers.

Discussion. Our results on the gradual retrieval and refinement of knowledge suggest an alternative perspective on how knowledge is encoded in language models. This viewpoint is different from recent work suggesting that models utilize a key-value memory structure, where facts are local to specific layers (Geva et al., 2022; Meng et al., 2022; Arad et al., 2023). Our results indicate that some information is distributed across layers, allowing for a gradual retrieval of knowledge rather than a retrieval at a particular point in the model. This aligns with earlier research proposing hierarchical representations in vision models (Zeiler and Fergus, 2014; Zhou et al., 2014; Bau et al., 2017).

6 Analyzing Model Failures

In this section, we delve into cases where the T2I model fails to generate images that align with the

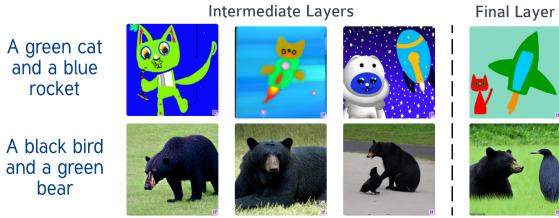


Figure 13: Examples of failure cases of T2I models (right). Using the DIFFUSION LENS (left) we can observe different patterns. In the first case (top row), the model is able to correctly generate each entity separately, but fails to combine them in the final layer. On the other hand (bottom row), the model is unable to generate a green bear in any of the intermediate representations.

input prompt. First, we investigate all failure cases in all experiments. Figure 14 shows the percentage of failures for each experiment that had over 10 failures. We split failures to two types: *complete failures* where no layer generated a correct image through DIFFUSION LENS, and cases where at least one layer generated a correct image, but the top layer led to a failure (*success then failure*).

Generally, the percentage of failure cases (total height of each bar) is low, from 10% to 25% for most categories. Prompts about two colored objects have a higher failure rate. Importantly, in many failure cases, the representations in earlier layers lead to a correct generation via our method. Notably, in simple prompts (relations and colored objects), about 80% of the failures had successful generations at earlier layers – see Figure 15 for an example. Once more constraints are imposed (two colored objects), we have a lower rate of early success. Finally, for knowledge-related tasks (famous people, unfamiliar animals), there are very few cases of early success turned to failure. Presumably, when the model fails, it is mostly because it does not encode the information at all.

Next, we zoom in on prompts describing two entities with different colors, as these prompts lead to the highest failure rate in our experiments. Examples for failure cases are shown in figure 13. Examining the final layer output images, the failures look similar: in both cases one or more entity was generated in the wrong color. However, using the DIFFUSION LENS, we reveal two different failure patterns: In the first example (“A green cat and a blue rocket”), both the blue rocket and the green cat are generated separately successfully in intermediate layers, while final output image fails

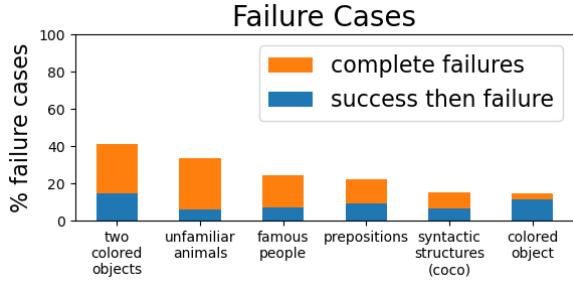


Figure 14: Many cases display successful generations from earlier layers before turning into failures.



Figure 15: DIFFUSION LENS reveals a correct image generation at a middle layer, while the final image fails to fully represent the prompt.

to combine them into a single image. This suggests that the failure stems from an unsuccessful combination of the two concepts. In the second example (“A black bird and a green bear”), the bear consistently appears black across all intermediate layers, signifying that the model struggles to generate a green bear throughout the encoding process of the text. A possible explanation is that black bear is a type of animal, which might mean that the phrase “black bear” is common in the training data, thus the appearance of the phrase “black” in the prompt biases the model towards generating black bears. This analysis reveals two different sources of failures that occur in compound prompts: either (1) the model fails at coupling a particular concept and color because it is biased towards another color, or that (2) the model can successfully couple each concept and color but fails to combine them.

We analyze how frequently each source of failure occurs, focusing on prompts that failed to generate a correct image from the final layer in at least 75% of cases. For each entity, we count the number of times it appeared in the correct color in at least one early layer. We find that for 40% of the failure cases in DF (70% in SD), at least one of the entities did not appear at all in images from earlier layers (type 1). The remaining set of failures had the correct color for each of the objects appear at some point in the computation (type 2).

7 Related Work

Interpreting language models. A wide range of work has analyzed language model internals. We briefly mention a few directions and refer to existing surveys (Belinkov and Glass, 2019; Rogers et al., 2020; Madsen et al., 2022; Ferrando et al., 2024). Probing classifier are used to analyze whether internal representations correlate with external properties (e.g., Ettinger et al., 2016; Hupkes et al., 2018). However, probing has various inherent flaws such as memorization (Belinkov, 2022), and requires costly annotations for fine-grained analysis like the visual characteristics of a specific animal species or person. Interventions in representations measure how they impact a model’s prediction (e.g., Vig et al., 2020; Elazar et al., 2021; Meng et al., 2022), and while they offer powerful insights, they are also challenging to design (Zhang and Nanda, 2023) and limited in scope. In contrast, the DIFFUSION LENS proposes a simple yet generic mechanism to visually interpret intermediate representations without requiring additional data or training, enabling exploration of fine-grained visual features.

Another influential approach is the Logit Lens (nostalgebraist, 2020), which projects intermediate representations of language models onto a probability distribution over the vocabulary space. The logit lens captures the internal computation of the language model, and the flow of information across modules (Geva et al., 2022; Katz and Belinkov, 2023; Pal et al., 2023). This line of work has focused on auto-regressive decoder language models. Inspired by this idea, we propose using the diffusion module in T2I pipelines to visualize intermediate prompt representations, revealing the text encoder’s computation process.

Interpreting vision-language models. Compared to unimodal models, research on interpretability in multimodal vision-language models is rather limited. Goh et al. (2021) found multi-modal neurons responding to specific concepts in CLIP (Radford et al., 2021) and Gandelsman et al. (2023) decomposed CLIP’s image representations into text-based characteristics.

Tang et al. (2023) were the first to propose a method to interpret T2I pipelines, by analyzing the influence of input words on generated images via cross-attention layers. Chefer et al. (2023b) decomposed textual concepts, focusing on the diffusion component. In contrast, our work investigates the

under-explored text encoder in T2I pipelines. Unlike previous methods, the DIFFUSION LENS reveals gradual processes within the model, not focusing only on the final output.

8 Discussion and Conclusion

We introduce the DIFFUSION LENS, a novel method to analyze language models within T2I pipelines. Our approach deconstructs the T2I pipeline by examining the output of each block within the text encoder, thereby providing a deeper insight into language-to-visual concept translation. We are the first, to our knowledge, to propose a method to interpret the text encoder and its internal computation process in the context of T2I models. Given that the text encoder is a crucial component of T2I models, enhancing its interpretability contributes to a deeper understanding of the entire generation process. We showcased the method’s potential by analyzing two open-source text encoders used in T2I pipeline across diverse topics.

Our work contributes to a growing body of work on analyzing how models process information across various components. The DIFFUSION LENS may have many potential applications as a first method of its kind, including similar applications to prior interpretability techniques such as improving model efficiency (Din et al., 2023; Dalvi et al., 2020) and tracing factual associations in language models, facilitating more accurate model editing methods (Meng et al., 2022; Arad et al., 2023).

Other future directions using the DIFFUSION LENS may aid in identifying points of failure in the computation process or remove undesired traits from early layers such as hallucinations, toxicity, or incorrect factual knowledge. Lastly, while we focused on entire blocks, our approach paves the way for visualizing individual sub-block components such as individual MLPs, attention heads, and residual connections.

Acknowledgements

This research was supported by the Israel Science Foundation (grant 448/20), an Azrieli Foundation Early Career Faculty Fellowship, and an AI Alignment grant from Open Philanthropy. HO is supported by the Apple AIML PhD fellowship. DA is supported by the Ariane de Rothschild Women Doctoral Program.

Limitations

While the DIFFUSION LENS provides a method to interpret the intermediate representations of the text encoder of T2I models, there are several limitations.

First, we are limited by the number of publicly available and open source T2I models and their corresponding text encoders. Extending our method to interpret other language models, whether or not they are used in T2I pipelines, offers a promising direction for future research.

Additionally, most of our experiments utilized automatically generated prompts, used to isolate and investigate specific effects. Such synthetic prompts are often less complex compared to prompts written by humans, and follow specific patterns. Although we experimented with a set of natural prompts, further exploration using a wider range of prompts could provide deeper insights into the behavior of text encoders in T2I models.

Lastly, the DIFFUSION LENS requires further annotation in order to derive large-scale conclusions. In this work, we relied on human and automatic annotation to answer questions on specific attributes of the generated images. This limitation stems from using images as the output of our method, however, we believe using images results in richer and more complex interpretations.

Ethics Statement

In this work, our primary objective is to enhance the transparency of text-to-image models. While not the focus of our analyses, the DIFFUSION LENS has the potential to unveil biases within these models. We anticipate that our work will contribute positively to the ongoing discourse on ethical practices in text-to-image models. At present, we do not foresee major ethical concerns arising from our methodology.

References

- Dana Arad, Hadas Orgad, and Yonatan Belinkov. 2023. ReFACT: Updating text-to-image models by editing the text encoder. *arXiv preprint arXiv:2306.00738*.
- David Bau, Bolei Zhou, Aditya Khosla, Aude Oliva, and Antonio Torralba. 2017. Network dissection: Quantifying interpretability of deep visual representations. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 6541–6549.
- Yonatan Belinkov. 2022. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219.
- Yonatan Belinkov and James Glass. 2019. [Analysis methods in neural language processing: A survey](#). *Transactions of the Association for Computational Linguistics*, 7:49–72.
- Nora Belrose, Zach Furman, Logan Smith, Danny Halawi, Igor Ostrovsky, Lev McKinney, Stella Biderman, and Jacob Steinhardt. 2023. Eliciting latent predictions from transformers with the tuned lens. *arXiv preprint arXiv:2303.08112*.
- Hila Chefer, Yuval Alaluf, Yael Vinker, Lior Wolf, and Daniel Cohen-Or. 2023a. Attend-and-excite: Attention-based semantic guidance for text-to-image diffusion models. *ACM Transactions on Graphics*, 42(4):1–10.
- Hila Chefer, Oran Lang, Mor Geva, Volodymyr Polosukhin, Assaf Shocher, Michal Irani, Inbar Mosseri, and Lior Wolf. 2023b. The hidden language of diffusion models. *arXiv preprint arXiv:2306.00966*.
- Fahim Dalvi, Hassan Sajjad, Nadir Durrani, and Yonatan Belinkov. 2020. [Analyzing redundancy in pretrained transformer models](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 4908–4926, Online. Association for Computational Linguistics.
- Alexander Yom Din, Taelin Karidi, Leshem Choshen, and Mor Geva. 2023. Jump to conclusions: Shortcutting transformers with linear transformations. *arXiv preprint arXiv:2303.09435*.
- Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. 2021. Amnesic probing: Behavioral explanation with amnesic counterfactuals. *Transactions of the Association for Computational Linguistics*, 9:160–175.
- Allyson Ettinger, Ahmed Elgohary, and Philip Resnik. 2016. [Probing for semantic evidence of composition by means of simple classification tasks](#). In *Proceedings of the 1st Workshop on Evaluating Vector-Space Representations for NLP*, pages 134–139, Berlin, Germany. Association for Computational Linguistics.
- Javier Ferrando, Gabriele Sarti, Arianna Bisazza, and Marta R Costa-jussà. 2024. A primer on the inner workings of transformer-based language models. *arXiv preprint arXiv:2405.00208*.
- Yossi Gandelsman, Alexei A Efros, and Jacob Steinhardt. 2023. Interpreting CLIP’s image representation via text-based decomposition. *arXiv preprint arXiv:2310.05916*.
- Mor Geva, Avi Caciularu, Kevin Wang, and Yoav Goldberg. 2022. [Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 30–45, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

- Gabriel Goh, Nick Cammarata, Chelsea Voss, Shan Carter, Michael Petrov, Ludwig Schubert, Alec Radford, and Chris Olah. 2021. Multimodal neurons in artificial neural networks. *Distill*, 6(3):e30.
- James Hampton. 2013. Conceptual combination 1. In *Knowledge Concepts and Categories*, pages 133–159. Psychology Press.
- Amir Hertz, Ron Mokady, Jay Tenenbaum, Kfir Aberman, Yael Pritch, and Daniel Cohen-Or. 2023. [Prompt-to-prompt image editing with cross-attention control](#). In *The Eleventh International Conference on Learning Representations, 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net.
- Dieuwke Hupkes, Sara Veldhoen, and Willem Zuidema. 2018. Visualisation and ‘diagnostic classifiers’ reveal how recurrent and recursive neural networks process hierarchical structure. *Journal of Artificial Intelligence Research*, 61:907–926.
- Gabriel Ilharco, Mitchell Wortsman, Ross Wightman, Cade Gordon, Nicholas Carlini, Rohan Taori, Achal Dave, Vaishaal Shankar, Hongseok Namkoong, John Miller, Hannaneh Hajishirzi, Ali Farhadi, and Ludwig Schmidt. 2021. [Openclip](#). If you use this software, please cite it as below.
- Shahar Katz and Yonatan Belinkov. 2023. Visit: Visualizing and interpreting the semantic information flow of transformers. *Findings of The 2023 Conference on Empirical Methods in Natural Language Processing*.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. 2023. BLIP-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. *arXiv preprint arXiv:2301.12597*.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, Lubomir Bourdev, Ross Girshick, James Hays, Pietro Perona, Deva Ramanan, C. Lawrence Zitnick, and Piotr Dollár. 2015. [Microsoft coco: Common objects in context](#).
- Ling ling Wu and Lawrence W. Barsalou. 2009. [Perceptual simulation in conceptual combination: Evidence from property generation](#). *Acta Psychologica*, 132(2):173–189. Spatial working memory and imagery: From eye movements to grounded cognition.
- Andreas Madsen, Siva Reddy, and Sarath Chandar. 2022. Post-hoc interpretability for neural NLP: A survey. *ACM Computing Surveys*, 55(8):1–42.
- Henry B Mann and Donald R Whitney. 1947. On a test of whether one of two random variables is stochastically larger than the other. *The annals of mathematical statistics*, pages 50–60.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. 2022. [Locating and editing factual associations in GPT](#). In *NeurIPS*.
- nostalgebraist. 2020. [Interpreting GPT: The logit lens](#). lesswrong, 2020.
- OpenAI. 2023. [GPT-4 technical report](#).
- Hadas Orgad, Bahjat Kawar, and Yonatan Belinkov. 2023. Editing implicit assumptions in text-to-image diffusion models. *arXiv preprint arXiv:2303.08084*.
- Koyena Pal, Jiuding Sun, Andrew Yuan, Byron Wallace, and David Bau. 2023. [Future lens: Anticipating subsequent tokens from a single hidden state](#). In *Proceedings of the 27th Conference on Computational Natural Language Learning*, pages 548–560, Singapore. Association for Computational Linguistics.
- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Kopf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. 2019. [Pytorch: An imperative style, high-performance deep learning library](#). In *Advances in Neural Information Processing Systems 32*, pages 8024–8035. Curran Associates, Inc.
- Peng Qi, Yuhao Zhang, Yuhui Zhang, Jason Bolton, and Christopher D Manning. 2020. Stanza: A python natural language processing toolkit for many human languages. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 101–108.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. 2021. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pages 8748–8763. PMLR.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. 2020. Exploring the limits of transfer learning with a unified text-to-text transformer. *The Journal of Machine Learning Research*, 21(1):5485–5551.
- Aditya Ramesh, Prafulla Dhariwal, Alex Nichol, Casey Chu, and Mark Chen. 2022. Hierarchical text-conditional image generation with CLIP latents. *arXiv preprint arXiv:2204.06125*, 1(2):3.
- Anna Rogers, Olga Kovaleva, and Anna Rumshisky. 2020. [A primer in BERTology: What we know about how BERT works](#). *Transactions of the Association for Computational Linguistics*, 8:842–866.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. 2022. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10684–10695.
- Chitwan Saharia, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily L Denton, Kamkar Ghazemipour, Raphael Gontijo Lopes, Burcu

- Karagol Ayan, Tim Salimans, et al. 2022. Photo-realistic text-to-image diffusion models with deep language understanding. *Advances in Neural Information Processing Systems*, 35:36479–36494.
- Dvir Samuel, Rami Ben-Ari, Simon Raviv, Nir Darshan, and Gal Chechik. 2024. Generating images of rare concepts using pre-trained diffusion models.
- Christoph Schuhmann, Romain Beaumont, Richard Vencu, Cade Gordon, Ross Wightman, Mehdi Cherti, Theo Coombes, Aarush Katta, Clayton Mullis, Mitchell Wortsman, et al. 2022. Laion-5b: An open large-scale dataset for training next generation image-text models. *Advances in Neural Information Processing Systems*, 35:25278–25294.
- Neil J Smelser, Paul B Baltes, et al. 2001. *International encyclopedia of the social & behavioral sciences*, volume 11. Elsevier Amsterdam.
- StabilityAI. 2023. Deepfloyd if.
- Raphael Tang, Linqing Liu, Akshat Pandey, Zhiying Jiang, Gefei Yang, Karun Kumar, Pontus Stenetorp, Jimmy Lin, and Ferhan Ture. 2023. What the DAAM: Interpreting stable diffusion using cross attention. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5644–5659, Toronto, Canada. Association for Computational Linguistics.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems*, 30.
- Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart Shieber. 2020. Investigating gender bias in language models using causal mediation analysis. In *Advances in Neural Information Processing Systems*.
- Patrick von Platen, Suraj Patil, Anton Lozhkov, Pedro Cuenca, Nathan Lambert, Kashif Rasul, Mishig Davaadorj, and Thomas Wolf. 2022. Diffusers: State-of-the-art diffusion models. <https://github.com/huggingface/diffusers>.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pieric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, et al. 2020. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 conference on empirical methods in natural language processing: system demonstrations*, pages 38–45.
- Yutaro Yamada, Yingtian Tang, and Ilker Yildirim. 2022. When are lemons purple? the concept association bias of CLIP. *arXiv preprint arXiv:2212.12043*.
- Matthew D. Zeiler and Rob Fergus. 2014. Visualizing and understanding convolutional networks. In *13th European Conference on Computer Vision, Zurich, Switzerland, September 6-12, 2014, Proceedings, Part I 13*, pages 818–833, Cham. Springer International Publishing.
- Fred Zhang and Neel Nanda. 2023. Towards best practices of activation patching in language models: Metrics and methods. *arXiv preprint arXiv:2309.16042*.
- Bolei Zhou, Agata Lapedriza, Jianxiong Xiao, Antonio Torralba, and Aude Oliva. 2014. Learning deep features for scene recognition using places database. *Advances in neural information processing systems*, 27.

A Additional Results

A.1 Prepositions

We explore prepositions in prompts. We investigate how prompts, including certain relations, affect the generation process. These prompts are complex, challenging the compositional understanding of the T2I model. In particular, we examine the prepositions "on" and "in". Figure 16 illustrates the percentage of images that correctly generated the concepts for tree categories: each of the objects alone and both objects with the correct relation between them. Our findings reveal that the emergence of each of the objects occurs at an early stage. However, both objects and their correct relation emerge only later in the text encoding.

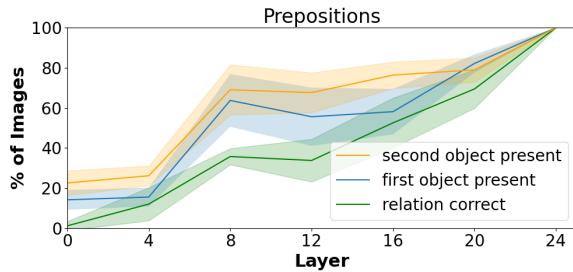


Figure 16: The proportion of images where either the objects, or objects with prepositions, were accurately represented.

A.2 Race between objects

Figure 18 presents examples of "race" between the objects in the prompts: one object appears first, and then disappears at a later layer to make room for the other object, before finally emerging again in the top layers.

A.3 Final layer norm necessity

In the DIFFUSION LENS process, we pass the output of block l through the last layer norm ln_f . However, we examine the option to bypass the ln_f layer and directly connect to the components of the diffusion model. As Figure 19 demonstrates, images generated without the final layer normalization are meaningless. The final layer norm thus plays a crucial role in generating meaningful images. It highlights the necessity of the ln_f layer within DIFFUSION LENS pipeline. A similar finding has been observed in the LogitLens (nostalgebraist, 2020) and TunedLens (Belrose et al., 2023).

B Annotation Process

The results in this paper rely on human annotators to determine the presence of different concepts in the generated images. We employed a team of ten professional full-time annotators using the Dataloop platform , in accordance with institutional regulations. The annotator teams was based in India, and were paid a rate of 8 USD per hour, in accordance with laws in India.

Each annotator received the instructions in Figure 20. The annotators were given the instruction to be liberal towards a positive answer. We manually validated each question, making sure the concepts in the question are not abstract (e.g., "beautiful"), and that the answer should be clear for each case. For each experiment, we duplicate 10% of the images, and ask an additional annotator the same questions, used to calculate inter annotator agreement. For experiments containing rare animals and celebrities, annotators were given reference images from google.

We provide our main results based on the human annotations. We chose to use human annotations since the existing automatic methods are limited. CLIP as an image classifier was shown to fail when required to explicitly bind attributes to objects (Ramesh et al., 2022; Yamada et al., 2022), and exploratory experiments we performed with BLIP (Li et al., 2023) showed similar issues.

We found a high agreement between GPT-4V (OpenAI, 2023) and the human annotators on most tasks and questions, as shown in Table 2. For one experiment – two colored objects – we found a high variance using the human annotations and thus extended it to further annotations using GPT4-V.

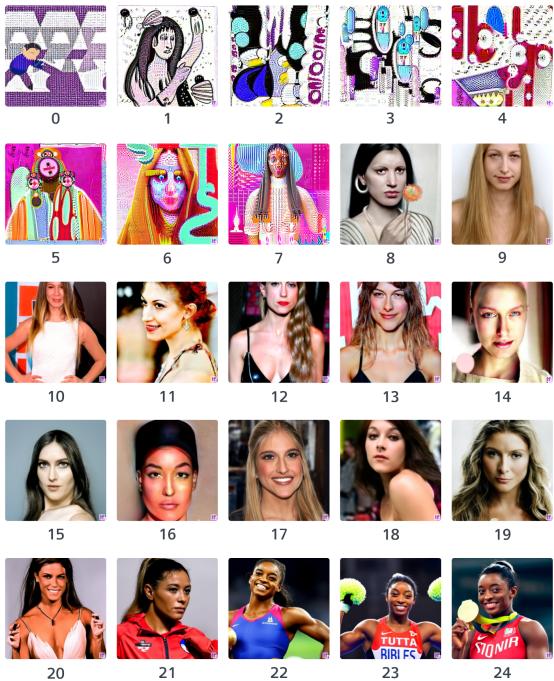
C Animals Experiment: Implementation Details

C.1 Animal classes used

To measure the gradual knowledge retrieval, one of the questions we ask in the experiment on unfamiliar animals is whether the image contains an animal of class X, where we vary X according to an informal, popular taxonomy that the specific animal belongs to. Note that although it does not faithfully represent the scientific view on the animals we generate, it is more suitable to observe a model that was trained on data that was taken from the wide internet.

To verify the distinction between familiar and unfamiliar animal species we preformed a Mann-

Simone Biles - Deep Floyd



Simone Biles - Stable Diffusion

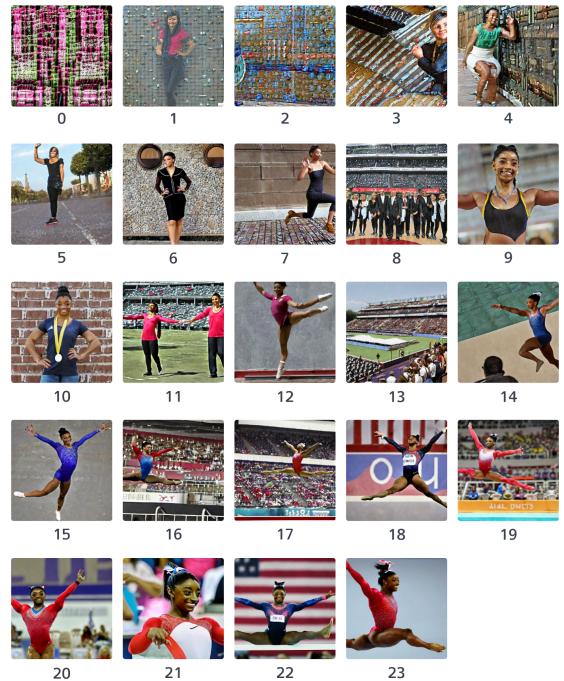


Figure 17: Example generations from all layers



Figure 18: A sequential “race” between two objects in the sentence, where one initially appears before the other, only to subsequently vanish and make room for the latter object.

Whitney U rank test (Mann and Whitney, 1947) on the frequencies of species names in the LAION2B-en dataset (Schuhmann et al., 2022), commonly used in the training process of T2I models which was computed by (Samuel et al., 2024). We found that the frequency of familiar species was greater than that of unfamiliar species with a confidence level of 95%.

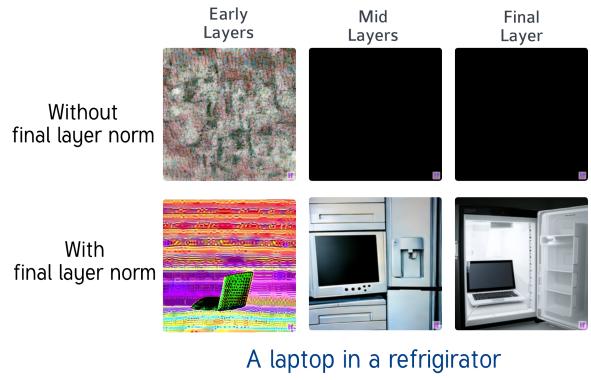


Figure 19: Example generations from DIFFUSION LENS with and without the final layer norm.

C.2 The full list of animals

Familiar animals: Beagle, German Shepherd, Labrador Retriever, Dachshund, Bulldog, Ragdoll, Kangaroo, Chicken, Owl, Eagle, Salmon, Catfish, Cod, Orca, Komodo dragon, King cobra, Platypus, Narwhal, Ostrich, cougar.

Unfamiliar animals: Aye-aye, Dik-dik, Tarsier, Gerenuk, Jerboa, Babirusa, Saola, Galago, Vervet, guppy, Celestial Pearl Danio, Herring, Pike, Wall-eye, Grebe, Spoonbill, Bee-eater, Taipan, ,Copper-

| Question type | Inter annotator agreements | | | Agreements with automatic annotations | | |
|-------------------------------------|----------------------------|-------|---------------|---------------------------------------|-------|---------------|
| | #annotations | f1 | cohen's kappa | #annotations | f1 | cohen's kappa |
| One object presence (out of 2) | 416 | 72.5% | 48.2% | 1381 | 80.6% | 63.8% |
| Relation correct | 208 | 73.7% | 61.4% | 1319 | 81.3% | 70.1% |
| One Color presence | 208 | 76.9% | 60.7% | 1671 | 85.3% | 85.9% |
| Familiar animals presence | 52 | 94.7% | 87.2% | 789 | 85.5% | 67.2% |
| Unfamiliar animals presence | 104 | 84.6% | 81.3% | 1019 | 84.3% | 72.4% |
| Unfamiliar animals class presence | 260 | 73.2% | 59.5% | 1012 | 91.2% | 81.3% |
| Syntactic structures correct (coco) | 357 | 80.6% | 69.7% | 2962 | 80.0% | 59.5% |

Table 2: A table of agreement between human annotators (left) and between human and automatic annotations averaged over both models. Overall, we see a high agreement between the human annotators and between the human and automatic annotations. For human agreement - the lowest Kappa score is for one object presence, probably due to the ambiguity in early layers, where there is a mix of both objects. For example in fig 5, second line, layer 12.

On this project, you will have to annotate sets of 50 images. For each set, you will have a yes or no question. The questions are written at the start of each task name. They end with a “?”. The latter part of the name is in “[]” and is not relevant for the questions. For convenience, we start the question with the statement itself, therefore “dog in the image?” means “Is there a dog in the image?” The questions vary from simple questions like “Is there a dog in the image?” to more complicated questions like “Is there a red bird on a green boat?”. The images are generated by AI, and might not be realistic. You should answer if the image might be interpreted as the question asks. Examples at the end of this file.

Figure 20: Annotation guidelines.

head, Anilius, Skink, Bearded Dragon, Ladybug, Scarab, Blue morpho, Cloudless sulphur, Giant anteater

D Implementation Details

We implemented our code using Pytorch (Paszke et al., 2019) and Huggingface libraries (Wolf et al., 2020; von Platen et al., 2022). For each experiment, we generated four images (different seeds) for each layer, and we report the standard deviation over the seeds in all plots. We use Stable Diffusion v2-1 (CreativeML Open RAIL-++-M License) (Rombach et al., 2022) and Deep Floyd (DeepFloyd-IF-License) (StabilityAI, 2023). We ran the experiments on the following GPUs: Nvidia A40, RTX 6000 Ada Generation, RTX A4000 and GeForce RTX 2080 Ti.

Our code is available in the supplementary material.

D.1 Dependency parsing implementation

We conducted a syntactic structure analysis using Stanza (Qi et al., 2020), a Python package. Stanza provides tools for obtaining parts of speech (POS) and syntactic structure dependency parse. To per-

form this analysis, we executed a Stanza pipeline designed for English. This pipeline returns the tokenized form, POS, lemmatization, and syntactic dependency parsing for a given prompt. We didn’t customize any additional parameters and utilized the default settings during the analysis.

E Results on Stable Diffusion

To complement the results in the main paper, we provide Figures 21–28 from Stable Diffusion.

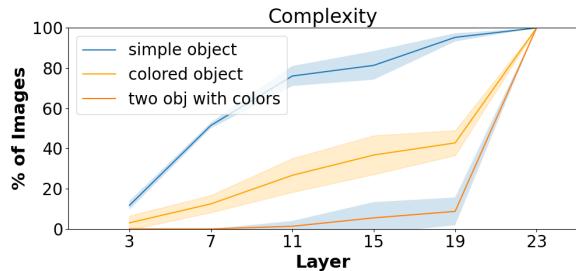


Figure 21: [Stable Diffusion] The percentage of images, from each category, for which the prompt matches the generated image, across different intermediate layers.

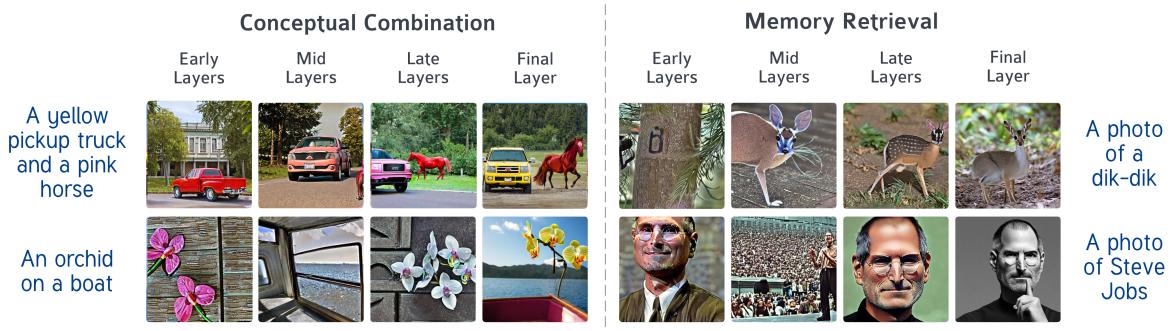


Figure 22: [Stable Diffusion] Insights gained from using DIFFUSION LENS. **Conceptual Combination** (left): early layers often act as a “bag of concepts”, lacking relational information which emerges in later layers. **Memory Retrieval** (right): concepts emerge early and gradually refine over layers.

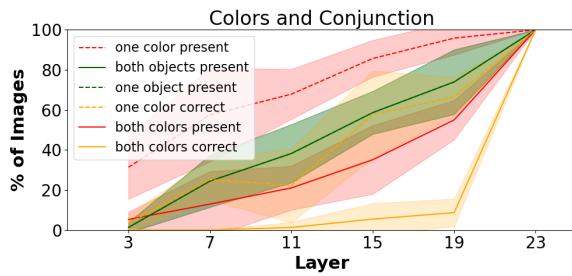


Figure 23: [Stable Diffusion] The proportion of images where either the object, the colors, or both were present, and where either the objects or the colors were accurately represented.

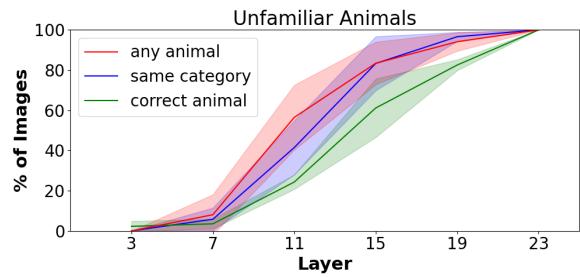


Figure 26: [Stable Diffusion] Subset of layers encoding different features in the process of unfamiliar animal generation.

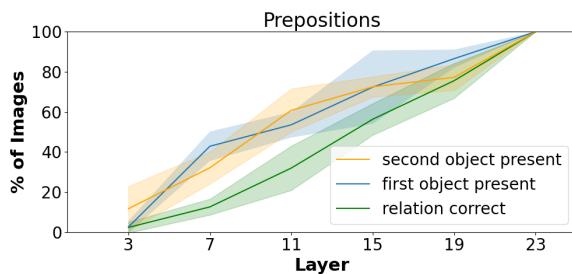


Figure 24: [Stable Diffusion] The proportion of images where either the objects, or objects with prepositions, were accurately represented.

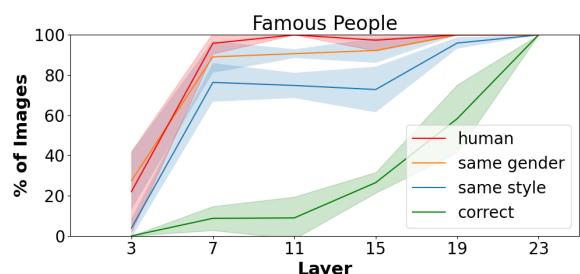


Figure 27: [Stable Diffusion] The distribution of feature granularity across layers in generated images.

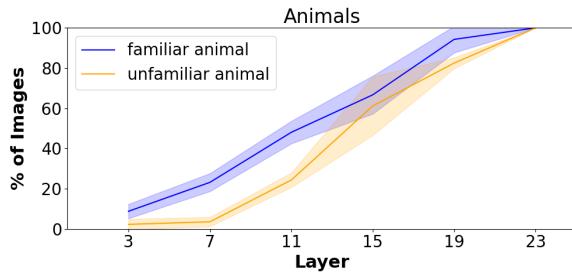


Figure 25: [Stable Diffusion] Familiar vs. unfamiliar animals across layers.

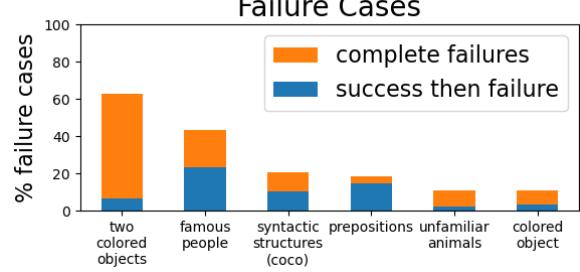


Figure 28: [Stable Diffusion] Many cases display successful generations from earlier layers before turning into failures.